

# Capstone Project

## Image classifier for the SVHN dataset

### Instructions

In this notebook, you will create a neural network that classifies real-world images digits. You will use concepts from throughout this course in building, training, testing, validating and saving your Tensorflow classifier model.

This project is peer-assessed. Within this notebook you will find instructions in each section for how to complete the project. Pay close attention to the instructions as the peer review will be carried out according to a grading rubric that checks key parts of the project instructions. Feel free to add extra cells into the notebook as required.

### How to submit

When you have completed the Capstone project notebook, you will submit a pdf of the notebook for peer review. First ensure that the notebook has been fully executed from beginning to end, and all of the cell outputs are visible. This is important, as the grading rubric depends on the reviewer being able to view the outputs of your notebook. Save the notebook as a pdf (File -> Download as -> PDF via LaTeX). You should then submit this pdf for review.

### Let's get started!

We'll start by running some imports, and loading the dataset. For this project you are free to make further imports throughout the notebook as you wish.

In [1]:

```
import tensorflow as tf
from scipy.io import loadmat
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Flatten, Conv2D, MaxPooling2D, Dropout, BatchNormalization
from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping
```

For the capstone project, you will use the [SVHN dataset](#). This is an image dataset of over 600,000 digit images in all, and is a harder dataset than MNIST as the numbers appear in the context of natural scene images. SVHN is obtained from house numbers in Google Street View images.

- Y. Netzer, T. Wang, A. Coates, A. Bissacco, B. Wu and A. Y. Ng. "Reading Digits in Natural Images with Unsupervised Feature Learning". NIPS Workshop on Deep Learning and Unsupervised Feature Learning, 2011.

Your goal is to develop an end-to-end workflow for building, training, validating, evaluating and saving a neural network that classifies a real-world image into one of ten classes.

In [2]:

```
# Run this cell to load the dataset

train = loadmat('data/train_32x32.mat')
test = loadmat('data/test_32x32.mat')
```

Both `train` and `test` are dictionaries with keys `x` and `y` for the input images and labels respectively.

## 1. Inspect and preprocess the dataset

- Extract the training and testing images and labels separately from the train and test dictionaries loaded for you.
- Select a random sample of images and corresponding labels from the dataset (at least 10), and display them in a figure.
- Convert the training and test images to grayscale by taking the average across all colour channels for each pixel. *Hint: retain the channel dimension, which will now have size 1.*
- Select a random sample of the grayscale images and corresponding labels from the dataset (at least 10), and display them in a figure.

In [3]:

```
X_train = train['X']
y_train = train['y']
X_test = test['X']
y_test = test['y']
```

In [5]:

```
import matplotlib.pyplot as plt
import numpy as np
fig, axs = plt.subplots(nrows=4, ncols=5, figsize=(10,8), subplot_kw={'xticks': [], 'yticks': []})
list_labels = []
for i in range(1, 21):
    rnd_num = np.random.randint(X_train.shape[3], size=1)
    img = X_train[:, :, :, rnd_num]
    fig.add_subplot(4, 5, i)
    plt.imshow(img.squeeze())
    plt.axis('off')
    list_labels.append(y_train[rnd_num])

for i, ax in enumerate(axs.reshape(-1)):
    title = str(list_labels[i]).strip("[]")
    ax.set_title(title)

plt.show()
```



In [6]:

```
X_train = X_train.mean(axis=2, dtype="float32").reshape(32, 32, 1, -1)
X_test = X_test.mean(axis=2, dtype="float32").reshape(32, 32, 1, -1)
print(X_train.shape)
print(X_test.shape)
```


```
(32, 32, 1, 73257)
(32, 32, 1, 26032)
```

In [7]:

```
fig, axs = plt.subplots(nrows =4, ncols=5, figsize=(10,8), subplot_kw={'xticks': [], 'yticks': []})
list_labels = []
for i in range(1, 21):
    rnd_num = np.random.randint(X_train.shape[3], size=1)
    img = X_train[:, :, :, rnd_num]
    fig.add_subplot(4, 5, i)
    plt.imshow(img.squeeze())
    plt.axis('off')
    list_labels.append(y_train[rnd_num])

for i, ax in enumerate(axs.reshape(-1)):
    title = str(list_labels[i]).strip("[]")
    ax.set_title(title)

plt.show()
```



## 2. MLP neural network classifier

- Build an MLP classifier model using the Sequential API. Your model should use only Flatten and Dense layers, with the final layer having a 10-way softmax output.
- You should design and build the model yourself. Feel free to experiment with different MLP architectures. *Hint: to achieve a reasonable accuracy you won't need to use more than 4 or 5 layers.*
- Print out the model summary (using the summary() method)
- Compile and train the model (we recommend a maximum of 30 epochs), making use of both training and validation sets during the training run.
- Your model should track at least one appropriate metric, and use at least two callbacks during training, one of which should be a ModelCheckpoint callback.
- As a guide, you should aim to achieve a final categorical cross entropy training loss of less than 1.0 (the

validation loss might be higher).

- Plot the learning curves for loss vs epoch and accuracy vs epoch for both training and validation sets.
- Compute and display the loss and accuracy of the trained model on the test set.

In [8]:

```
X_train = X_train.T.reshape(-1, 32, 32, 1)
X_train = np.swapaxes(X_train, axis1=1, axis2=2)
X_test = X_test.T.reshape(-1, 32, 32, 1)
X_test = np.swapaxes(X_test, axis1=1, axis2=2)
```

```
y_train = y_train.squeeze()
y_test = y_test.squeeze()
print(X_train.shape)
print(X_test.shape)
```

```
(73257, 32, 32, 1)
(26032, 32, 32, 1)
```

In [9]:

```
X_train = X_train / 255.0
X_test = X_test / 255.0
```

In [10]:

```
list_labels = np.unique(y_train)
list_labels
```

Out[10]:

```
array([ 1,  2,  3,  4,  5,  6,  7,  8,  9, 10], dtype=uint8)
```

In [12]:

```
from tensorflow.keras.utils import to_categorical
y_train_truth_value = to_categorical(y_train - 1, num_classes=10)
y_test_truth_value = to_categorical(y_test - 1, num_classes=10)
```

In [14]:

```
w1 = tf.keras.initializers.RandomNormal(mean=0.0, stddev=0.05)
model = Sequential([
    Flatten(input_shape=(32,32,1)),
    Dense(512, activation="relu", kernel_initializer=w1),
    Dense(256, activation="relu", kernel_initializer=w1),
    Dense(256, activation="relu", kernel_initializer=w1),
    Dense(10, activation="softmax")
])
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
flatten (Flatten)	(None, 1024)	0
dense (Dense)	(None, 512)	524800
dense_1 (Dense)	(None, 256)	131328
dense_2 (Dense)	(None, 256)	65792
dense_3 (Dense)	(None, 10)	2570
=====		
Total params: 724,490		
Trainable params: 724,490		
Non-trainable params: 0		

```
In [25]:
```

```
model.compile(optimizer='adam', loss="categorical_crossentropy", metrics=["accuracy"])
```

```
In [26]:
```

```
checkpoint_best = ModelCheckpoint(filepath = 'model_checkpoints_best/checkpoint',
                                   save_weights_only=True,
                                   save_freq='epoch',
                                   monitor='loss',
                                   mode='min',
                                   save_best_only=True,
                                   verbose=1)
early = EarlyStopping(monitor='loss', patience=3, verbose=1)
```

```
In [16]:
```

```
history = model.fit(X_train, y_train_truth_value, batch_size=128, epochs=30, validation_split=0.15, callbacks=[checkpoint_best, early])
```

Train on 62268 samples, validate on 10989 samples

Epoch 1/30

62208/62268 [=====>.] - ETA: 0s - loss: 1.9417 - accuracy: 0.3058

ETA: 1s - loss: 1.9553

Epoch 00001: loss improved from inf to 1.94128, saving model to model\_checkpoints\_best/checkpoint

62268/62268 [=====] - 38s 618us/sample - loss: 1.9413 - accuracy : 0.3060 - val\_loss: 1.4821 - val\_accuracy: 0.4977

Epoch 2/30

62080/62268 [=====>.] - ETA: 0s - loss: 1.2895 - accuracy: 0.5783

Epoch 00002: loss improved from 1.94128 to 1.28850, saving model to model\_checkpoints\_best/checkpoint

62268/62268 [=====] - 36s 579us/sample - loss: 1.2885 - accuracy : 0.5786 - val\_loss: 1.2391 - val\_accuracy: 0.5981

Epoch 3/30

62080/62268 [=====>.] - ETA: 0s - loss: 1.0743 - accuracy: 0.6601

Epoch 00003: loss improved from 1.28850 to 1.07424, saving model to model\_checkpoints\_best/checkpoint

62268/62268 [=====] - 36s 575us/sample - loss: 1.0742 - accuracy : 0.6601 - val\_loss: 1.0500 - val\_accuracy: 0.6639

Epoch 4/30

62208/62268 [=====>.] - ETA: 0s - loss: 0.9654 - accuracy: 0.6984

Epoch 00004: loss improved from 1.07424 to 0.96564, saving model to model\_checkpoints\_best/checkpoint

62268/62268 [=====] - 36s 575us/sample - loss: 0.9656 - accuracy : 0.6984 - val\_loss: 0.9319 - val\_accuracy: 0.7063

Epoch 5/30

62208/62268 [=====>.] - ETA: 0s - loss: 0.8921 - accuracy: 0.7223

Epoch 00005: loss improved from 0.96564 to 0.89224, saving model to model\_checkpoints\_best/checkpoint

62268/62268 [=====] - 36s 573us/sample - loss: 0.8922 - accuracy : 0.7223 - val\_loss: 0.9014 - val\_accuracy: 0.7155

Epoch 6/30

62208/62268 [=====>.] - ETA: 0s - loss: 0.8322 - accuracy: 0.7413

Epoch 00006: loss improved from 0.89224 to 0.83236, saving model to model\_checkpoints\_best/checkpoint

62268/62268 [=====] - 35s 568us/sample - loss: 0.8324 - accuracy : 0.7413 - val\_loss: 0.8277 - val\_accuracy: 0.7448

Epoch 7/30

62080/62268 [=====>.] - ETA: 0s - loss: 0.7903 - accuracy: 0.7539

Epoch 00007: loss improved from 0.83236 to 0.79002, saving model to model\_checkpoints\_best/checkpoint

62268/62268 [=====] - 36s 576us/sample - loss: 0.7900 - accuracy : 0.7540 - val\_loss: 0.7928 - val\_accuracy: 0.7512

Epoch 8/30

62208/62268 [=====>.] - ETA: 0s - loss: 0.7522 - accuracy: 0.7647

Epoch 00008: loss improved from 0.79002 to 0.75207, saving model to model\_checkpoints\_best/checkpoint

62268/62268 [=====] - 36s 574us/sample - loss: 0.7521 - accuracy : 0.7647 - val\_loss: 0.8161 - val\_accuracy: 0.7484

Epoch 9/30

62208/62268 [=====>.] - ETA: 0s - loss: 0.7162 - accuracy: 0.7769

Epoch 00009: loss improved from 0.75207 to 0.71611, saving model to model\_checkpoints\_best/checkpoint

```
Epoch 0009: loss improved from 0.73207 to 0.71611, saving model to model_checkpoints_bes
t/checkpoint
62268/62268 [=====] - 35s 569us/sample - loss: 0.7161 - accuracy
: 0.7769 - val_loss: 0.7110 - val_accuracy: 0.7772
Epoch 10/30
62208/62268 [=====>.] - ETA: 0s - loss: 0.6897 - accuracy: 0.7845
Epoch 0010: loss improved from 0.71611 to 0.68981, saving model to model_checkpoints_bes
t/checkpoint
62268/62268 [=====] - 35s 569us/sample - loss: 0.6898 - accuracy
: 0.7845 - val_loss: 0.7288 - val_accuracy: 0.7751
Epoch 11/30
62080/62268 [=====>.] - ETA: 0s - loss: 0.6627 - accuracy: 0.7923
Epoch 0011: loss improved from 0.68981 to 0.66303, saving model to model_checkpoints_bes
t/checkpoint
62268/62268 [=====] - 35s 570us/sample - loss: 0.6630 - accuracy
: 0.7923 - val_loss: 0.7110 - val_accuracy: 0.7758
Epoch 12/30
62208/62268 [=====>.] - ETA: 0s - loss: 0.6349 - accuracy: 0.8008
Epoch 0012: loss improved from 0.66303 to 0.63511, saving model to model_checkpoints_bes
t/checkpoint
62268/62268 [=====] - 35s 568us/sample - loss: 0.6351 - accuracy
: 0.8008 - val_loss: 0.7343 - val_accuracy: 0.7642
Epoch 13/30
62208/62268 [=====>.] - ETA: 0s - loss: 0.6204 - accuracy: 0.8053
Epoch 0013: loss improved from 0.63511 to 0.62040, saving model to model_checkpoints_bes
t/checkpoint
62268/62268 [=====] - 36s 574us/sample - loss: 0.6204 - accuracy
: 0.8053 - val_loss: 0.7048 - val_accuracy: 0.7730
Epoch 14/30
62080/62268 [=====>.] - ETA: 0s - loss: 0.5935 - accuracy: 0.8140
Epoch 0014: loss improved from 0.62040 to 0.59347, saving model to model_checkpoints_bes
t/checkpoint
62268/62268 [=====] - 36s 574us/sample - loss: 0.5935 - accuracy
: 0.8140 - val_loss: 0.6360 - val_accuracy: 0.8057
Epoch 15/30
62080/62268 [=====>.] - ETA: 0s - loss: 0.5806 - accuracy: 0.8156
Epoch 0015: loss improved from 0.59347 to 0.58028, saving model to model_checkpoints_bes
t/checkpoint
62268/62268 [=====] - 35s 569us/sample - loss: 0.5803 - accuracy
: 0.8158 - val_loss: 0.6564 - val_accuracy: 0.7963
Epoch 16/30
62208/62268 [=====>.] - ETA: 0s - loss: 0.5777 - accuracy: 0.8163
Epoch 0016: loss improved from 0.58028 to 0.57765, saving model to model_checkpoints_bes
t/checkpoint
62268/62268 [=====] - 36s 576us/sample - loss: 0.5776 - accuracy
: 0.8163 - val_loss: 0.6298 - val_accuracy: 0.8067
Epoch 17/30
62208/62268 [=====>.] - ETA: 0s - loss: 0.5512 - accuracy: 0.8247
Epoch 0017: loss improved from 0.57765 to 0.55123, saving model to model_checkpoints_bes
t/checkpoint
62268/62268 [=====] - 35s 570us/sample - loss: 0.5512 - accuracy
: 0.8247 - val_loss: 0.6186 - val_accuracy: 0.8097
Epoch 18/30
62208/62268 [=====>.] - ETA: 0s - loss: 0.5348 - accuracy: 0.8306
Epoch 0018: loss improved from 0.55123 to 0.53498, saving model to model_checkpoints_bes
t/checkpoint
62268/62268 [=====] - 35s 569us/sample - loss: 0.5350 - accuracy
: 0.8306 - val_loss: 0.6201 - val_accuracy: 0.8118
Epoch 19/30
62208/62268 [=====>.] - ETA: 0s - loss: 0.5200 - accuracy: 0.8342
Epoch 0019: loss improved from 0.53498 to 0.51997, saving model to model_checkpoints_bes
t/checkpoint
62268/62268 [=====] - 36s 572us/sample - loss: 0.5200 - accuracy
: 0.8342 - val_loss: 0.6536 - val_accuracy: 0.7989
Epoch 20/30
62208/62268 [=====>.] - ETA: 0s - loss: 0.5085 - accuracy: 0.8367
ETA: 0s - loss: 0.5082 - accuracy: 0.
Epoch 0020: loss improved from 0.51997 to 0.50875, saving model to model_checkpoints_bes
t/checkpoint
62268/62268 [=====] - 36s 570us/sample - loss: 0.5087 - accuracy
: 0.8367 - val_loss: 0.6224 - val_accuracy: 0.8072
Epoch 21/30
62208/62268 [=====>.] - ETA: 0s - loss: 0.4978 - accuracy: 0.8405
```

```

62208/62268 [=====>.] - ETA: 0s - loss: 0.4970 - accuracy: 0.8403
Epoch 00021: loss improved from 0.50875 to 0.49771, saving model to model_checkpoints_best/checkpoint
62268/62268 [=====] - 35s 569us/sample - loss: 0.4977 - accuracy
: 0.8406 - val_loss: 0.6256 - val_accuracy: 0.8067
Epoch 22/30
62208/62268 [=====>.] - ETA: 0s - loss: 0.4757 - accuracy: 0.8473
Epoch 00022: loss improved from 0.49771 to 0.47544, saving model to model_checkpoints_best/checkpoint
62268/62268 [=====] - 36s 574us/sample - loss: 0.4754 - accuracy
: 0.8474 - val_loss: 0.6280 - val_accuracy: 0.8094
Epoch 23/30
62208/62268 [=====>.] - ETA: 0s - loss: 0.4730 - accuracy: 0.8496
Epoch 00023: loss improved from 0.47544 to 0.47293, saving model to model_checkpoints_best/checkpoint
62268/62268 [=====] - 36s 575us/sample - loss: 0.4729 - accuracy
: 0.8497 - val_loss: 0.6136 - val_accuracy: 0.8187
Epoch 24/30
62208/62268 [=====>.] - ETA: 0s - loss: 0.4607 - accuracy: 0.8535
ETA: 0s - loss: 0.4615
Epoch 00024: loss improved from 0.47293 to 0.46080, saving model to model_checkpoints_best/checkpoint
62268/62268 [=====] - 36s 571us/sample - loss: 0.4608 - accuracy
: 0.8535 - val_loss: 0.5900 - val_accuracy: 0.8226
Epoch 25/30
62208/62268 [=====>.] - ETA: 0s - loss: 0.4446 - accuracy: 0.8580
ETA: 0s - loss: 0.4448 - accuracy: 0.85
Epoch 00025: loss improved from 0.46080 to 0.44449, saving model to model_checkpoints_best/checkpoint
62268/62268 [=====] - 36s 572us/sample - loss: 0.4445 - accuracy
: 0.8580 - val_loss: 0.6006 - val_accuracy: 0.8220
Epoch 26/30
62208/62268 [=====>.] - ETA: 0s - loss: 0.4351 - accuracy: 0.8611
Epoch 00026: loss improved from 0.44449 to 0.43495, saving model to model_checkpoints_best/checkpoint
62268/62268 [=====] - 36s 571us/sample - loss: 0.4349 - accuracy
: 0.8612 - val_loss: 0.5984 - val_accuracy: 0.8224
Epoch 27/30
62208/62268 [=====>.] - ETA: 0s - loss: 0.4245 - accuracy: 0.8629
Epoch 00027: loss improved from 0.43495 to 0.42474, saving model to model_checkpoints_best/checkpoint
62268/62268 [=====] - 35s 564us/sample - loss: 0.4247 - accuracy
: 0.8628 - val_loss: 0.6652 - val_accuracy: 0.8023
Epoch 28/30
62208/62268 [=====>.] - ETA: 0s - loss: 0.4187 - accuracy: 0.8664
Epoch 00028: loss improved from 0.42474 to 0.41875, saving model to model_checkpoints_best/checkpoint
62268/62268 [=====] - 35s 565us/sample - loss: 0.4188 - accuracy
: 0.8664 - val_loss: 0.6199 - val_accuracy: 0.8146
Epoch 29/30
62080/62268 [=====>.] - ETA: 0s - loss: 0.4166 - accuracy: 0.8661
Epoch 00029: loss improved from 0.41875 to 0.41651, saving model to model_checkpoints_best/checkpoint
62268/62268 [=====] - 35s 562us/sample - loss: 0.4165 - accuracy
: 0.8661 - val_loss: 0.5737 - val_accuracy: 0.8320
Epoch 30/30
62208/62268 [=====>.] - ETA: 0s - loss: 0.3955 - accuracy: 0.8730
Epoch 00030: loss improved from 0.41651 to 0.39560, saving model to model_checkpoints_best/checkpoint
62268/62268 [=====] - 35s 557us/sample - loss: 0.3956 - accuracy
: 0.8730 - val_loss: 0.6015 - val_accuracy: 0.8244


```

In [17]:

```


fig, axs = plt.subplots(figsize=(6,6))
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.xlabel('epoch')
plt.ylabel('loss')
plt.legend(['train', 'test'], loc="upper left")
plt.show()

```



In [18]:

```
fig, axs = plt.subplots(figsize=(6,6))
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.xlabel('epoch')
plt.ylabel('accuracy')
plt.legend(['train', 'test'], loc="upper left")
plt.show()
```



In [30]:

```
test_loss, test_accuracy = model.evaluate(X_test, y_test_truth_value, verbose=False)
print("Loss on test samples:", np.around(test_loss, 2))
print("Accuracy on test samples:", np.around(test_accuracy*100,2), '%')
```

Loss on test samples: 0.74  
Accuracy on test samples: 80.2 %

### 3. CNN neural network classifier

- **Build a CNN classifier model using the Sequential API. Your model should use the Conv2D, MaxPool2D, BatchNormalization, Flatten, Dense and Dropout layers. The final layer should again have a 10-way softmax output.**





```
mode='max',
save_best_only=True,
verbose=1)
early_CNN = EarlyStopping(monitor='val_accuracy', patience=3, verbose=1)
```

In [18]:

```
history_CNN = CNN.fit(X_train, y_train_truth_value, batch_size=128, epochs=30, validation_split=0.15, callbacks=[checkpoint_best_CNN, early_CNN])
```

Train on 62268 samples, validate on 10989 samples

Epoch 1/30

62208/62268 [=====>.] - ETA: 0s - loss: 1.1054 - accuracy: 0.6344

Epoch 00001: val\_accuracy improved from -inf to 0.80599, saving model to model\_checkpoint\_s\_best\_CNN/checkpoint

62268/62268 [=====] - 124s 2ms/sample - loss: 1.1048 - accuracy: 0.6346 - val\_loss: 0.9149 - val\_accuracy: 0.8060

Epoch 2/30

62208/62268 [=====>.] - ETA: 0s - loss: 0.5780 - accuracy: 0.8189

Epoch 00002: val\_accuracy improved from 0.80599 to 0.84039, saving model to model\_checkpoint\_s\_best\_CNN/checkpoint

62268/62268 [=====] - 123s 2ms/sample - loss: 0.5778 - accuracy: 0.8189 - val\_loss: 0.5111 - val\_accuracy: 0.8404

Epoch 3/30

62208/62268 [=====>.] - ETA: 0s - loss: 0.4857 - accuracy: 0.8467

Epoch 00003: val\_accuracy improved from 0.84039 to 0.86514, saving model to model\_checkpoint\_s\_best\_CNN/checkpoint

62268/62268 [=====] - 123s 2ms/sample - loss: 0.4855 - accuracy: 0.8467 - val\_loss: 0.4425 - val\_accuracy: 0.8651

Epoch 4/30

62208/62268 [=====>.] - ETA: 0s - loss: 0.4363 - accuracy: 0.8640

Epoch 00004: val\_accuracy did not improve from 0.86514

62268/62268 [=====] - 120s 2ms/sample - loss: 0.4364 - accuracy: 0.8640 - val\_loss: 0.5412 - val\_accuracy: 0.8432

Epoch 5/30

62208/62268 [=====>.] - ETA: 0s - loss: 0.4016 - accuracy: 0.8751

Epoch 00005: val\_accuracy improved from 0.86514 to 0.88452, saving model to model\_checkpoint\_s\_best\_CNN/checkpoint

62268/62268 [=====] - 120s 2ms/sample - loss: 0.4016 - accuracy: 0.8751 - val\_loss: 0.3819 - val\_accuracy: 0.8845

Epoch 6/30

62208/62268 [=====>.] - ETA: 0s - loss: 0.3751 - accuracy: 0.8828

Epoch 00006: val\_accuracy improved from 0.88452 to 0.88716, saving model to model\_checkpoint\_s\_best\_CNN/checkpoint

62268/62268 [=====] - 121s 2ms/sample - loss: 0.3750 - accuracy: 0.8828 - val\_loss: 0.3804 - val\_accuracy: 0.8872

Epoch 7/30

62208/62268 [=====>.] - ETA: 0s - loss: 0.3481 - accuracy: 0.8924

Epoch 00007: val\_accuracy did not improve from 0.88716

62268/62268 [=====] - 121s 2ms/sample - loss: 0.3481 - accuracy: 0.8924 - val\_loss: 0.4019 - val\_accuracy: 0.8819

Epoch 8/30

62208/62268 [=====>.] - ETA: 0s - loss: 0.3315 - accuracy: 0.8962

Epoch 00008: val\_accuracy improved from 0.88716 to 0.89753, saving model to model\_checkpoint\_s\_best\_CNN/checkpoint

62268/62268 [=====] - 121s 2ms/sample - loss: 0.3314 - accuracy: 0.8962 - val\_loss: 0.3561 - val\_accuracy: 0.8975

Epoch 9/30

62208/62268 [=====>.] - ETA: 0s - loss: 0.3146 - accuracy: 0.9012

Epoch 00009: val\_accuracy did not improve from 0.89753

62268/62268 [=====] - 121s 2ms/sample - loss: 0.3148 - accuracy: 0.9012 - val\_loss: 0.4286 - val\_accuracy: 0.8817

Epoch 10/30

62208/62268 [=====>.] - ETA: 0s - loss: 0.3019 - accuracy: 0.9067

Epoch 00010: val\_accuracy improved from 0.89753 to 0.90172, saving model to model\_checkpoint\_s\_best\_CNN/checkpoint

62268/62268 [=====] - 120s 2ms/sample - loss: 0.3018 - accuracy: 0.9067 - val\_loss: 0.3377 - val\_accuracy: 0.9017

Epoch 11/30

62208/62268 [=====>.] - ETA: 0s - loss: 0.2904 - accuracy: 0.9089


Epoch 00011: val\_accuracy did not improve from 0.90172

62268/62268 [=====] - 121s 2ms/sample - loss: 0.2903 - accuracy:

```
0.9089 - val_loss: 0.3548 - val_accuracy: 0.8957
Epoch 12/30
62208/62268 [=====>.] - ETA: 0s - loss: 0.2677 - accuracy: 0.9158
Epoch 00013: val_accuracy did not improve from 0.90172
62268/62268 [=====] - 121s 2ms/sample - loss: 0.2678 - accuracy:
0.9157 - val_loss: 0.3822 - val_accuracy: 0.8925
Epoch 00013: early stopping
```


In [19]:

```
fig, axs = plt.subplots()
plt.plot(history_CNN.history['loss'])
plt.plot(history_CNN.history['val_loss'])
plt.xlabel('epoch')
plt.ylabel('loss')
plt.legend(['train', 'test'], loc="upper left")
plt.show()
```



In [20]:

```
fig, axs = plt.subplots(figsize=(6,6))
plt.plot(history_CNN.history['accuracy'])
plt.plot(history_CNN.history['val_accuracy'])
plt.xlabel('epoch')
plt.ylabel('loss')
plt.legend(['train', 'test'], loc="upper left")
plt.show()
```



In [22]:

```
test_loss_CNN, test_accuracy_CNN = CNN.evaluate(X_test, y_test_truth_value, verbose=False)
```

```
)
print("Loss on test samples:", np.around(test_loss_CNN, 2))
print("Accuracy on test samples:", np.around(test_accuracy_CNN*100,2), '%')
```

Loss on test samples: 0.42  
Accuracy on test samples: 88.4 %

## 4. Get model predictions

- Load the best weights for the MLP and CNN models that you saved during the training run.
- Randomly select 5 images and corresponding labels from the test set and display the images with their labels.
- Alongside the image and label, show each model's predictive distribution as a bar chart, and the final model prediction given by the label with maximum probability.

In [23]:

```
model.load_weights('model_checkpoints_best/checkpoint')
CNN.load_weights('model_checkpoints_best_CNN/checkpoint')
```

Out[23]:

<tensorflow.python.training.tracking.util.CheckpointLoadStatus at 0x7fbe4b9fc3c8>

In [26]:

```
images = []
labels = []
predictions_MLP = []
predictions_CNN = []





for i in range(0, 5):
    rnd_num = np.random.randint(X_test.shape[0], size=1)
    img = X_test[rnd_num, :, :, :]
    images.append(img)
    labels.append(str(y_test[rnd_num]).strip('[]'))
    predictions_MLP.append(model.predict(img))
    predictions_CNN.append(CNN.predict(img))
```




In [27]:

```
fig, axs = plt.subplots(nrows=5, ncols=3, figsize=(20,20))
fig.subplots_adjust(hspace=0.5, wspace=0.2)




x = np.arange(1,11)

for i, (image, label, mlp_prediction, cnn_prediction) in enumerate(zip(images, labels, p
redictions_MLP, predictions_CNN)):
    axs[i, 0].imshow(np.squeeze(image))
    axs[i, 0].set_title(label)
    axs[i, 0].axis('off')
    axs[i, 1].bar(x, mlp_prediction.reshape(10,))
    axs[i, 1].set_title('MLP Classifier Predictions')
    axs[i, 1].set_xticks(x)
    axs[i, 2].bar(x, cnn_prediction.reshape(10,))
    axs[i, 2].set_title('CNN Classifier Predictions')
    axs[i, 2].set_xticks(x)
```











2



4



7



In [ ]:

In [ ]:

In [ ]:

In [ ]: